



EasyBuild: building software with ease

a gently introduction & recent developments

Kenneth Hoste

HPC-UGent, Ghent University, Belgium

kenneth.hoste@ugent.be

http://users.ugent.be/~kehoste/EasyBuild-intro-TACC_20151111.pdf

EasyBuild/Lmod hackathon @ TACC, Austin (TX), US – Nov 11th 2015

HPC-UGent in a nutshell



<http://www.ugent.be/hpc>

- HPC team at central IT dept. of Ghent University (Belgium)
- 9 team members: 1 manager, ~3 user support, ~5 sysadmin
- 7 Tier2 clusters + one Tier1 (8.5k cores), >1k servers in total
- ~1.5k user accounts, across all scientific domains
- tasks: hardware, system administration, user support/training, ...

- member of Flemish Supercomputer Centre (VSC)
virtual centre, collaboration between Flemish university associations



“Please install this on the HPC?”

In the context of high performance computing, *building from source* should be preferred, when possible (if sources are available).

This allows for controlling used compilers and libraries, optimizing the software for the specific system architecture (e.g., AVX, network), etc.

Installing (lots of) *scientific* software is typically:

- error-prone, trial-and-error
- tedious, hard to get right
- repetitive & boring (well...)
- time-consuming (hours, days, even weeks)
- frustrating (“*Pandora’s box*”)
- sometimes simply not worth the effort...



Common issues with scientific software

Researchers focus on the *science* behind the software they implement, and care little about tools, build procedure, portability, . . .

Scientists are not software developers or sysadmins (nor should they be).

"If we would know what we are doing, it wouldn't be called 'research'."

This results in:

- use of non-standard build tools (or broken ones)
- incomplete build procedure, e.g., no configure or install step
- interactive installation scripts
- hardcoded parameters (compilers, libraries, paths, . . .)
- poor/outdated/missing/incorrect documentation
- dependency (version) hell

Prime example I: WRF

Weather Research and Forecasting Model (<http://www.wrf-model.org>)
(one of the top 5 applications on Blue Waters)

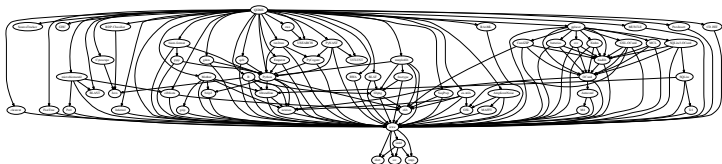
- dozen dependencies: netCDF (C, Fortran), HDF5, tcsh, JasPer, ...
- known issues in last release are (only) documented on website
no patch file provided, infrequent bugfix releases
- interactive 'configure' script :(
- resulting `configure.wrf` needs work:
fix hardcoded settings (compilers, libraries, ...), tweaking of options
- custom 'compile' script (wraps around 'make')
building in parallel is broken without fixing the Makefile
- no actual installation step

Wouldn't it be nice to build & install WRF with a single command?

http://easybuild.readthedocs.org/en/latest/Typical_workflow_example_with_WRF.html

Prime example II: QIIME

QIIME: Quantitative Insights Into Microbial Ecology (<http://qiime.org/>)



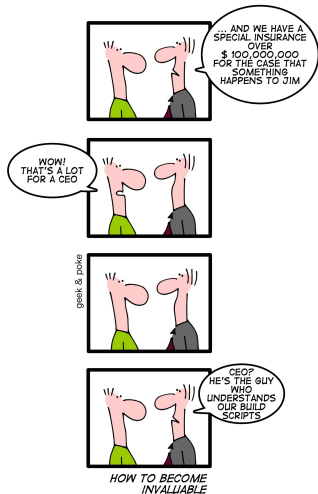
- scientific research domain: bioinformatics ...
- 59 dependencies in total (*without* compiler toolchain), some optional
 - depends on Haskell (GHC), Java, Python, R, Perl, OCaml, ...
 - several deps use a non-standard build procedure (in various degrees)
- very picky about dependency versions (e.g., *must* be Python v2.7.3)
- took us several weeks to get it installed (like we wanted)...
- ... **now we can (re)build/install it all with a single command!**

(disclaimer: support for QIIME not included yet in latest EasyBuild release)

Houston, we have a problem

Installation of scientific software is a *tremendous* problem for HPC sites all around the world.

- huge burden on HPC user support teams
- researchers lose lots of time (waiting)
- sites typically resort to in-house scripting
- very little collaboration among HPC sites :(



What about existing tools?

Existing tools are not well suited to scientific software and HPC systems.

- package managers: **yum** (RPMs), **apt-get** (.deb), ...
- **Homebrew** (Mac OS X), <http://brew.sh/>
- **Linuxbrew**, <http://brew.sh/linuxbrew/>
- **Portage** (Gentoo), <http://wiki.gentoo.org/wiki/Project:Portage>
- **pkgsrc** (NetBSD & (a lot) more), <http://pkgsrc.org/>
- **Nix**, <http://nixos.org/>

Common problems:

- usually poor support for multiple versions/builds existing side-by-side
- not flexible enough to deal with idiosyncrasies of scientific software
- hard to maintain (bash, heavy copy-pasting, ...)
- little support for scientific software, other compilers (not GCC), ...

EasyBuild: building software with ease



<http://hpcugent.github.io/easybuild/>

- tool for building/installing (scientific) software on HPC systems
- collection of Python packages and modules
- in-house since 2009, open-source (GPLv2) since Nov 2012
- thriving community: actively contributing, driving development
- new release every 6–8 weeks (latest: EasyBuild v2.4.0, Nov 10th 2015)
next release: planned before end of 2015 (v2.5.0)
- supports over 730 different software packages
including CP2K, GAMESS-US, GROMACS, NAMD, NWChem,
OpenFOAM, PETSc, QuantumESPRESSO, WRF, WPS, ...
- well documented: <http://easybuild.readthedocs.org>

Projects similar to EasyBuild

- **Spack** (LLNL) - <http://scalability-llnl.github.io/spack/>
- **Maali** (Pawsey) - <https://github.com/chrisbpawsey/maali/>
- **Smithy** (NICS, ORNL) - <http://anthonydigirolamo.github.io/smithy/>

Major differences with EasyBuild:

- slightly different approach
- smaller community
- fewer supported software packages
- less flexibility

Most (all?) are interested in switching to/merging with EasyBuild.

EasyBuild: system requirements

- main target platform (for now) is Linux x86_64 (HPC systems)
 - Red Hat-based systems (Scientific Linux, CentOS, RHEL, ...)
 - also other Linux distros: Debian, Ubuntu, OpenSUSE, SLES, ...
 - also (kind of) works on OS X
 - *no* Windows support (and none planned)
 - experimental support for Cray systems since EasyBuild v2.1.0
 - support for Linux@POWER systems is being looked into (TAMU)
- Python v2.6.x or more recent v2.x (no Python 3 support yet)
- a modules tool:
 - latest release of Tcl/C environment modules (version 3.2.10);
 - or one of the Tcl-only versions of environment modules;
 - or a recent version of *Lmod* (5.6.3 or more recent)
- (a system C/C++ compiler, to get started)

EasyBuild: feature highlights

- fully **autonomously** building and installing (scientific) software
 - automatic dependency resolution
 - automatic generation of module files (Tcl or Lua syntax)
- thorough **logging** of executed build/install procedure
- **archiving** of build specifications ('*easyconfig files*')
- highly **configurable**, via config files/environment/command line
- **dynamically extendable** with additional *easyblocks*, *toolchains*, etc.
- support for **custom module naming schemes** (incl. hierarchical)
- **comprehensively tested**: lots of unit tests, regression testing, ...
- actively developed, **collaboration** between various HPC sites
- worldwide **community**

'Quick' demo for the impatient

```
eb HPL-2.0-goolf-1.4.10-no-OFED.eb --robot
```

- **downloads** all required sources (best effort)
- **builds/installs** *goolf* toolchain (be patient) + HPL on top of it
goolf: GCC, OpenMPI, LAPACK, OpenBLAS, FFTW, ScaLAPACK
- **generates module file** for each installed software package
- default: source/build/install dir in `$HOME/.local/easybuild`
can be easily changed by configuring EasyBuild differently

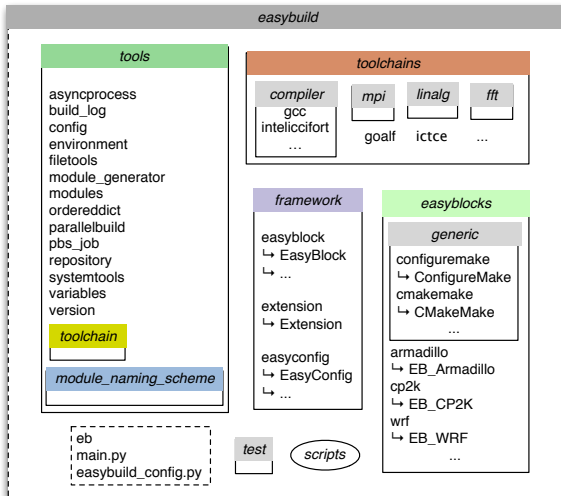
EasyBuild: high-level design overview

- EasyBuild *framework*
 - core of EasyBuild: Python modules & packages
 - provides supporting functionality for building and installing software
- *easyblock*
 - a Python module, 'plugin' for the EasyBuild framework
 - implements a (generic) software build/install procedure
- *easyconfig* file
 - build specification: software name/version, compiler toolchain, etc.
- compiler *toolchain*
 - compilers with accompanying libraries (MPI, BLAS/LAPACK, etc.)

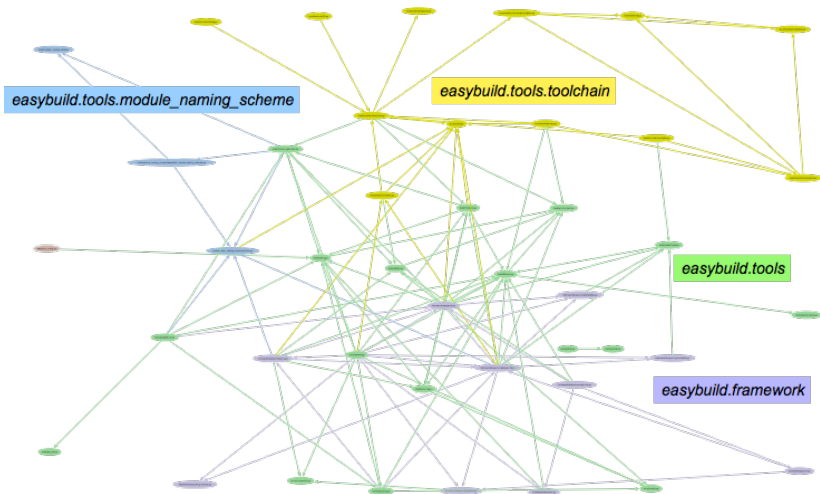
Putting it all together

The EasyBuild *framework* leverages *easyblocks* to automatically build and install (scientific) software using a particular *compiler toolchain*, as specified by one or more *easyconfig files*.

EasyBuild: high-level design overview

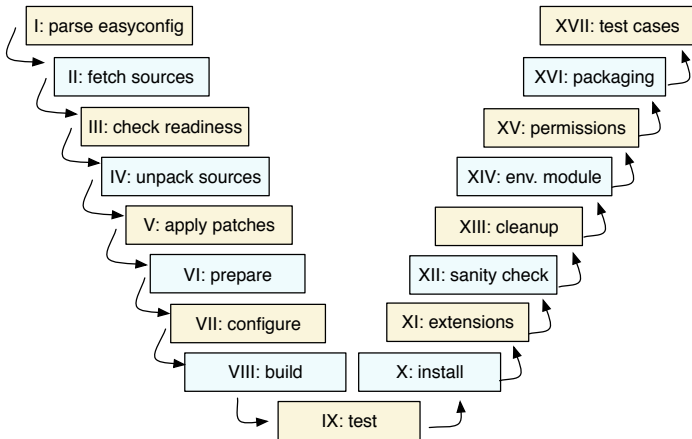


EasyBuild: high-level design overview



Step-wise install procedure

build and install procedure as implemented by EasyBuild



most of these steps can be customised if required,
via *easyconfig* parameters or a *custom easyblock*

EasyBuild: statistics

EasyBuild v2.4.0 (Nov'15)

- ~ 23,000 LoC in framework (17 Python packages, 153 Python modules)
 - + ~ 5,000 LoC in vsc-base (option parsing/logging)
 - + ~ 11,000 LoC more in unit tests
 - ⇒ ~ 39,000 LoC in total
- 182 easyblocks in total (~ 16,000 Loc)
 - 154 software-specific easyblocks
 - 28 generic easyblocks
- 784 different software packages supported (incl. toolchains & bundles)
 - bio: 172, tools: 104, vis: 82, lib: 78, devel 68, math: 49,
 - data: 50, toolchain: 38, chem: 32, lang: 28, mpi: 23,
 - numlib: 22, perf: 21, system: 19, cae: 16, compiler: 12, phys: 6
- 4,596 easyconfig files: different versions/variants, toolchains, ...

Installing EasyBuild

<http://easybuild.readthedocs.org/en/latest/Installation.html>

Install EasyBuild using the bootstrap script (*highly recommended*):

```
$ curl -O https://raw.githubusercontent.com/hpcugent/easybuild-framework/  
develop/easybuild/scripts/bootstrap_eb.py  
$ python bootstrap_eb.py /tmp/$USER # replace with your prefix  
$ module use /tmp/$USER/modules/all  
$ module load EasyBuild
```

Update EasyBuild with ... EasyBuild!

```
$ module load EasyBuild/2.3.0  
$ eb EasyBuild-2.4.0.eb  
$ module swap EasyBuild EasyBuild/2.4.0
```

Configuring EasyBuild

<http://easybuild.readthedocs.org/en/latest/Configuration.html>

By default, EasyBuild will (ab)use `$HOME/.local/easybuild`.

You should configure EasyBuild to your preferences, via:

- *configuration file(s)*: key-value lines, text files (e.g., `prefix=/tmp`)
- *environment variables* (e.g., `$EASYBUILD_PREFIX` set to `/tmp`)
- *command line parameters* (e.g., `--prefix=/tmp`)

Consistency across these options is guaranteed (see `eb --help | tail`).

Priority among different options: cmdline, env vars, config file.

For example:

- `--prefix` overrules `$EASYBUILD_PREFIX`
- `$EASYBUILD_PREFIX` overrules `prefix` in configuration file

Basic usage

http://easybuild.readthedocs.org/en/latest/Using_the_EasyBuild_command_line.html

- specify software name/version and toolchain to 'eb' command
- commonly via name(s) of easyconfig file(s):

```
eb GCC-4.9.2.eb Clang-3.6.0-GCC-4.9.2.eb
```

- or directory name(s) providing set(s) of easyconfig files:

```
eb $HOME/myeasyconfigs
```

- or via command line options:

```
eb --software-name=GCC
```

- --robot/-r: dependency resolution, --debug/-d: debug logging

```
eb WRF-3.6.1-intel-2015a-dmpar.eb -dr
```

Workflow example

http://easybuild.readthedocs.org/en/latest/Typical_workflow_example_with_WRF.html

- getting help, overview of options:

```
eb --help
```

- searching for easyconfigs:

```
eb --search HPL    # long output (full paths)
eb -S HPL          # short output
```

- overview of required/available modules:

```
eb HPL-2.1-foss-2015a.eb --dry-run  # long output
eb HPL-2.1-foss-2015a.eb -D        # short output
```

- robot build, enable debug logging, (also) log to stdout:

```
eb HPL-2.1-foss-2015a.eb --debug -lr
```

Other common command line options

- install to different installation prefix:

```
eb HPL-2.1-foss-2015a.eb --installpath=/tmp/$USER
```

- (try to) build & install different software version:

```
eb HPL-2.1-foss-2015a.eb --try-software-version=2.0
```

- (try to) build & install with a different toolchain (version):

```
eb HPL-2.0-foss-2014b.eb --try-toolchain-version=2015a -r
```

```
eb HPL-2.1-foss-2015a.eb --try-toolchain=intel,2015a -r
```

- grab (specific) easyconfigs from a GitHub pull request:

```
eb --from-pr 1239 OpenMPI-1.8.4-GCC-4.9.2.eb --robot
```

EasyBuild community

- ~150 subscribers to the EasyBuild mailing list
- ~20 active members on the #easybuild IRC channel
- bi-weekly conf calls: recent developments, discussing problems, ...
- users/contributors at HPC sites and companies around the world
incl. Flemish Supercomputer Centre sites, Jülich Supercomputer Centre, Univ. of Basel, OHRI (Canada), Univ. of Auckland, Bayer AG, Texas A&M, IMP/IMBA (Austria), Univ. of Luxembourg, Cyprus Institute, University of Colorado Boulder, CSCS (Switzerland), ...
- *"Getting Scientific Software Installed"* BoF sessions at ISC/SC
- 9 'hackathon' workshops (2-3 days) since Aug'12
 - Ghent (Belgium), Luxembourg, Nicosia (Cyprus), Jülich (Germany), Vienna (Austria), Basel (Switzerland), Espoo (Finland)
 - **10th hackathon at TACC (Austin, Texas), Nov'15, before SC15**
- **1st EasyBuild User Meeting: Jan 27-29 2016, Ghent**

HUST'14 paper

Modern Scientific Software Management Using EasyBuild and Lmod

Markus Geimer (JSC)

Kenneth Hoste (HPC-UGent)

Robert McLay (TACC)

http://hpcugent.github.io/easybuild/files/hust14_paper.pdf

- paper at HPC User Support Tools workshop (HUST'14 @ SC14)
- explains basics of module tools, EasyBuild and Lmod
- highlights issues with current approaches in software installation
- advocates use of a hierarchical module naming scheme
- presents EasyBuild and Lmod as adequate tools for software management on HPC systems

HUST'15 during SC15 in Austin: <http://hust15.github.io/>

EasyBuild: recent developments

- `eb --extended-dry-run`, `eb -x` [see docs]
- support for easyconfig files in YAML syntax (`.yeb`) **[WIP, see docs]**
 - included in EasyBuild v2.4.0 under `--experimental`
- other experimental features [see docs]
 - Cray support (WIP)
 - packaging software installed with EasyBuild (using FPM)
 - using minimal (sub)toolchains for dependencies

EasyBuild: future work

- further extend documentation: generic easyblocks, easyblocks API
- **stabilize Cray support**
- **support for more Lmod-specific features**
 - module families
 - module properties & sticky modules
- **making all easyblocks “dry run-aware”**
- command line support for contributing easyconfigs
'eb --new-pr' and 'eb --update-pr'
- support for rpath-style linking of libraries
- **extend support for .yeb easyconfigs to ‘fat’ easyconfigs**
- **redesign module generation, to provide full control over generated module files**
- join forces with Spack (LLNL)?

Do you want to know more?

- EasyBuild website: <http://hpcugent.github.io/easybuild>
- EasyBuild documentation: <http://easybuild.readthedocs.org>
- stable EasyBuild releases: <http://pypi.python.org/pypi/easybuild>
 - EasyBuild framework: <http://pypi.python.org/pypi/easybuild-framework>
 - easyblocks: <http://pypi.python.org/pypi/easybuild-easyblocks>
 - easyconfigs <http://pypi.python.org/pypi/easybuild-easyconfigs>
- source repositories on GitHub
 - EasyBuild meta package + docs: <https://github.com/hpcugent/easybuild>
 - EasyBuild framework: <https://github.com/hpcugent/easybuild-framework>
 - easyblocks: <https://github.com/hpcugent/easybuild-easyblocks>
 - easyconfigs: <https://github.com/hpcugent/easybuild-easyconfigs>
- EasyBuild mailing list: easybuild@lists.ugent.be
<https://lists.ugent.be/www/subscribe/easybuild>
- Twitter: @easy_build
- IRC: #easybuild on chat.freenode.net