

XALT: User Environment Tracking

Robert McLay, Mark Fahey, Reuben Budiardja, Sandra Sweat

The Texas Advanced Computing Center, Argonne National Labs, NICS

Nov. 17, 2015



XALT: What runs on the system

- A NSF Funded project: PI: Mark Fahey and Robert McLay
- A Census of what programs and libraries are run
- Running at TACC, NICS, U. Florida, KAUST, ...
- Integrates with TACC-Stats.

Design Goals

- Be extremely light-weight
- Provide provenance data: How?
- How many use a library or application?
- Collect Data into a Database for analysis.

Design: Linker

- The linker (ld) wrapper intercepts the user link line.
 - A shell script wrapper, ld which uses python scripts
 - Generate assembly code: key-value pairs
 - Capture tracemap output from ld
 - Transmit collected data in *.json format

Design: Launcher

- Program Launcher: mpirun, aprun, ibrun ...
 - A shell script wrapper is called which uses python scripts
 - Find Executable by parsing command
 - Collect executable info, shared libraries, env.
 - Transmit collected data in *.json format
- In future: This will be handled by the linker phase

Design: Transmission to DB

- File: collect nightly
- Syslog: Use Syslog filtering
- Direct to DB.

Lmod to XALT connection

- Lmod spider walks entire module tree.
- Can build A Reverse Map from paths to modules
- Can map program & libraries to modules.
- /opt/apps/i15/mv2_2_1/phdf5/1.8.14/lib/libhdf5.so.9 \Rightarrow
phdf5/1.8.14(intel/15.02:mvapich2/2.1)

Lmod: Priority Path

- Fixed Job Launcher: ibrun, aprun
- Variable Launchers: mpirun, mpiexec
- Priority Path:
`prepend_path{"PATH", "/opt/apps/xalt/1.0/bin",
priority=100}`

Database Changes (I)

- Tables sizes in XALT:

Table	Size in MB
join_run_env	199603.00
join_run_object	9388.00
join_link_object	5013.00
xalt_run	4613.00
xalt_object	4175.00
xalt_link	814.00

- join_run_env has 2.1 billion rows

Database Changes (II)

- Environment variables are important.
- But mainly for reproducing results
- Not SQL tests (mostly)

Database Changes (III): New Design

- Store complete env \Rightarrow compressed json blob
- Filter Env's with Accept Test followed by Reject Test
- Instead of 250 vars per job \Rightarrow 20 to 30.

Protecting XALT (I): UTF8 Characters

- Linux supports UTF8 Characters in file names, env. vars.
- Python supports UTF8 if you know what you are doing.
- Switch XALT to use `cursor.execute(query, (job_id, user, ...))`
- Where `query="INSERT INTO table VALUE(%s,%s)"`
- This prevent SQL injection: "johnny drop tables;"
- Also supports UTF8 characters.

Protecting XALT (II): PYTHONHOME,...

- Four Ways: LD_LIBRARY_PATH, PATH, PYTHONPATH, PYTHONHOME
- Solution: LD_LIBRARY_PATH="@ld_lib_path@" PATH=@python@ -E python-script ...
- Everything that depends on PATH must be hard coded
- basename \Rightarrow /bin/basename
- Unique install for each operating system.
- Programs move around.

Using XALT Data

- Targetted Outreach: Who will be affected
- Largemem Queue Overuse
- XALT and TACC-Stats

Publishing XALT Data

- Student Sandra Sweat
- Sanitized Data
- Community Codes Reported: Vasp*, WRF*, OpenFOAM*,
- users names : U012354, Charge Accounts: A12345
- Unique mapping, Added Field of Science

Tracking Non-mpi jobs (I)

- Originally we tracked only MPI Jobs
- By hijacking mpirun etc.
- Now we can use ELF binary format to track jobs

ELF Binary Format Trick

```
void myinit(int argc, char **argv)
{
    /* ... */
}
void myfini()
{
    /* ... */
}
__attribute__((section(".init_array"))) typeof(myinit) *__init = myinit;
__attribute__((section(".fini_array"))) typeof(myfini) *__fini = myfini;
```

Using the ELF Binary Format Trick

- This C code is compiled and linked in through the hijacked linker
- It can also be used with LD_PRELOAD
- We are only testing linking in directly

Downsides

- Currently, we only track task 0 jobs.
- MPMD programs will only record the Task 0 job.
- We also lose the ability to capture return exit status

Upsides (I)

- Can now track all executables period.
- Can now track “launcher” jobs.

Upsides (II)

- Do not need to write/maintain a parser for ibrun, mpirun ...
- Do not need to correctly jump over certain executables:
 - OK: `ibrun tacc_affinity user_program`
 - Not O.K: `ibrun env -u foo user_program`

Challenges (I)

- With both LD_PRELOAD and init.o linked in. \Rightarrow double records
- Do not want to track mv, cp, etc
- Only want to track some executables on compute nodes
- Do not want to get overwhelmed by the data.

Challenges (II)

- Still under construction
- None of these issues are show-stoppers
- Full Release in the New Year.

XALT Demo

- Show modules hierarchy
- `ml -raw show xalt`
- `type -a ld,mpirun`
- Build programs
- Run tests
- Run utf8 program
- Show database results

Conclusion

- Lmod:
 - Source: github.com/TACC/lmod.git, lmod.sf.net
 - Documentation: lmod.readthedocs.org
- XALT:
 - Source: github.com/Fahey-McLay/xalt.git, xalt.sf.net
 - Documentation: [doc/*.pdf](#)